

Homework 7

2020-10-07

Question 10.1

Using the same crime data set `uscrime.txt` as in Questions 8.2 and 9.1, find the best model you can using

- (a) a regression tree model, and
- (b) a random forest model.

In R, you can use the `tree` package or the `rpart` package, and the `randomForest` package. For each model, describe one or two qualitative takeaways you get from analyzing the results (i.e., don't just stop when you have a good model, but interpret it too).

Answer:

- a) Using regression tree model:

```
rm(list = ls())
uscrime <- read.delim("uscrime.txt")
```

Read the data.

Fitting the regression tree function to the crime data.


```
Regression Tree Model
library(tree)
uscrime_tree <- tree(Crime~., data = uscrime)
summary(uscrime_tree)

##
## Regression tree:
## tree(formula = Crime ~ ., data = uscrime)
## Variables actually used in tree construction:
## [1] "Po1" "Pop" "LF" "NW"
## Number of terminal nodes: 7
## Residual mean deviance: 47390 = 1896000 / 40
## Distribution of residuals:
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## -573.900 -98.300  -1.545    0.000 110.600  490.100
```

Deviance: This is a quality of fit statistic that is a generalization of sum of squared residuals

Tree Visualization:

```
plot(uscrime_tree)
text(uscrime_tree)
```



```
uscrime_tree$frame
```


##	var	n	dev	yval	splits.cutleft	splits.cutright
## 1	Po1	47	6880927.66	905.0851	<7.65	>7.65
## 2	Pop	23	779243.48	669.6087	<22.5	>22.5
## 4	LF	12	243811.00	550.5000	<0.5675	>0.5675
## 8	<leaf>	7	48518.86	466.8571		
## 9	<leaf>	5	77757.20	667.6000		
## 5	<leaf>	11	179470.73	799.5455		
## 3	NW	24	3604162.50	1130.7500	<7.65	>7.65
## 6	Pop	10	557574.90	886.9000	<21.5	>21.5
## 12	<leaf>	5	146390.80	1049.2000		
## 13	<leaf>	5	147771.20	724.6000		
## 7	Po1	14	2027224.93	1304.9286	<9.65	>9.65
## 14	<leaf>	6	170828.00	1041.0000		
## 15	<leaf>	8	1124984.88	1502.8750		

```
uscrime_tree$where
```

```
## 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25
26
## 6 13 4 13 9 10 12 13 6 5 13 6 6 5 6 12 4 13 10 13 6 4 12 9 5
13
## 27 28 29 30 31 32 33 34 35 36 37 38 39 40 41 42 43 44 45 46 47
## 4 12 13 6 4 12 6 9 10 9 6 5 6 12 5 4 6 10 4 10 9
```

#Manually compute R2. Is this a good measure of the quality of fit?
#Notice that we can only use averages of each leaf to make
#predictions.

```
yhat_tree <- predict(uscrime_tree)
plot(yhat_tree, uscrime$Crime)
```



Let us prune the tree leaves and find the quality of fit of our model.

```
prune.tree(uscrime_tree)$size
```

```
## [1] 7 6 5 4 3 2 1
```

```
prune.tree(uscrime_tree)$dev
```

```
## [1] 1895722 2013257 2276670 2632631 3364043 4383406 6880928
```

```
set.seed(42)
```

The tree has 7 leaves and observe that the deviance is highest for the leaf 1 while lowest deviance is for leaf7.

Run the k fold cross validation to find the deviance or number of misclassification as a function of cost complexity parameter $k = 10$, FUN is the function to do pruning.


```
cv.tree(object = uscrime_tree, FUN = prune.tree, k = 10)

## $size
## [1] 7 6 5 4 3 2 1
##
## $dev
## [1] 7986688 7986688 7885099 7920849 7707536 7035494 8536032
##
## $k
## [1]      -Inf  117534.9  263412.9  355961.8  731412.1 1019362.7 2497521.7
##
## $method
## [1] "deviance"
##
## attr(,"class")
## [1] "prune"          "tree.sequence"
```

we can do cross validation.

Example of manually pruning a tree in which we choose to only have 4 leaves
Tree with only 4 leaves:

```
uscrime_tree_prune4 <- prune.tree(uscrime_tree,best = 4)
plot(uscrime_tree_prune4)
text(uscrime_tree_prune4)
```



```
summary(uscrime_tree_prune4)
```

```
##
```

```
## Regression tree:
```

```
## snip.tree(tree = uscrime_tree, nodes = c(6L, 2L))
```

```
## Variables actually used in tree construction:
```

```
## [1] "Po1" "NW"
```

```
## Number of terminal nodes: 4
```

```
## Residual mean deviance: 61220 = 2633000 / 43
```

```
## Distribution of residuals:
```

```
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
```

```
## -573.90 -152.60   35.39    0.00  158.90  490.10
```

```
crimeTree4Predict <- predict(uscrime_tree_prune4, data = uscrime[,1:15])
```

```
RSSofTree4 <- sum((crimeTree4Predict - uscrime[,16])^2)
```

```
TSS <- sum((uscrime[,16] - mean(uscrime[,16]))^2)
```

```
R2ofTree4 <- 1 - RSSofTree4/TSS
```

```
R2ofTree4
```

```
## [1] 0.6174017
```


```
The R Squared error for Tree with 4 leaves is 0.6174017
```

```
Prune the tree to have 5 leaves
```

```
uscrime_tree_prune5 <- prune.tree(uscrime_tree,best = 5)
```

```
plot(uscrime_tree_prune5)
```

```
text(uscrime_tree_prune5)
```



```

summary(uscrime_tree_prune5)


##
## Regression tree:
## snip.tree(tree = uscrime_tree, nodes = c(4L, 6L))
## Variables actually used in tree construction:
## [1] "Po1" "Pop" "NW"
## Number of terminal nodes: 5
## Residual mean deviance: 54210 = 2277000 / 42
## Distribution of residuals:
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## -573.9 -107.5   15.5     0.0  122.8   490.1

crimeTree5Predict <- predict(uscrime_tree_prune5, data = uscrime[,1:15])
RSSofTree5 <- sum((crimeTree5Predict - uscrime[,16])^2)
R2ofTree5 <- 1 - RSSofTree5 / TSS
R2ofTree5

## [1] 0.6691333

uscrime_tree_prune6 <- prune.tree(uscrime_tree,best = 6)
plot(uscrime_tree_prune6)
text(uscrime_tree_prune6)

```



```
summary(uscrime_tree_prune6)
```

```
##
## Regression tree:
## snip.tree(tree = uscrime_tree, nodes = 4L)
## Variables actually used in tree construction:
## [1] "Po1" "Pop" "NW"
## Number of terminal nodes: 6
## Residual mean deviance: 49100 = 2013000 / 41
## Distribution of residuals:
##      Min. 1st Qu.  Median      Mean 3rd Qu.     Max.
## -573.900 -99.520  -1.545    0.000  122.800  490.100

crimeTree6Predict <- predict(uscrime_tree_prune6, data = uscrime[,1:15])
RSSofTree6 <- sum((crimeTree6Predict - uscrime[,16])^2)
R2ofTree6 <- 1 - RSSofTree6 / TSS
R2ofTree6

## [1] 0.7074149
```

Analysis :

I built 4 models with different number of leaves. Below are the results.
 The RSquared error for tree with 7 leaves = 0.47390
 The RSquared error for tree with 6 leaves = 0.7074149
 The RSquared error for tree with 5 leaves = 0.6691333

The Rsquared error for tree with 4 leaves = 0.6174017

The model of the tree with 4 leaves has a lower R Squared error and so looks like the best among the built models.

Q10.1b : Using Random Forests

Grow Random trees and set the number of predictors that you want to consider at each split. Using number of predictors / 3 is a good number to choose. Since we have 15 predictors, $15/3 \sim 4$ would be a good value for the number of predictors at each split.

```
library(randomForest)

## randomForest 4.6-14

## Type rfNews() to see new features/changes/bug fixes.

set.seed(42)
num_pred <- 4
uscrime_rf <- randomForest(Crime~.,data = uscrime,mtry = num_pred,importance
= TRUE , ntree = 500)
uscrime_rf

##
## Call:
## randomForest(formula = Crime ~ ., data = uscrime, mtry = num_pred,
importance = TRUE, ntree = 500)
##
##           Type of random forest: regression
##           Number of trees: 500
## No. of variables tried at each split: 4
##
##           Mean of squared residuals: 83636.77
##           % Var explained: 42.87

crime_rf_predict <- predict(uscrime_rf, data=uscrime[, -16])
RSS <- sum((crime_rf_predict - uscrime[,16])^2)
R2 <- 1 - RSS/TSS
R2

## [1] 0.4287212
```

The RSquared error for mtry = 4 is 0.4287212

```
num_pred5 <- 5
uscrime_rf5 <- randomForest(Crime~.,data = uscrime,mtry =
num_pred5,importance = TRUE , ntree = 500)
uscrime_rf5

##
## Call:
```



```
## randomForest(formula = Crime ~ ., data = uscrime, mtry = num_pred5,
importance = TRUE, ntree = 500)
##           Type of random forest: regression
##           Number of trees: 500
## No. of variables tried at each split: 5
##
##           Mean of squared residuals: 88034.11
##           % Var explained: 39.87
```

```
crime_rf_predict5 <- predict(uscrime_rf5, data=uscrime[, -16])
RSS5 <- sum((crime_rf_predict5 - uscrime[, 16])^2)
R2_pre5 <- 1 - RSS5/TSS
R2_pre5
```

```
## [1] 0.3986853
```

The RSquared error for mtry = 5 is 0.3986953

For mtry = 5 has the lowest RSquared than the model with mtry = 4.

```
importance(uscrime_rf)
```

```
##           %IncMSE IncNodePurity
## M           2.4984854      200566.40
## So           1.3802135       33881.59
## Ed           4.8378328      198601.72
## Po1          9.7354718     1076933.25
## Po2         10.6715396     1268930.03
## LF           0.6449124       311872.13
## M.F          1.1555044      239897.22
## Pop          2.1893155      379760.15
## NW           8.7310286      542658.76
## U1           2.6422460      145760.60
## U2           1.6754487      190587.49
## Wealth       3.2683848      626353.30
## Ineq         2.1162044      238557.90
## Prob         8.6884908      812217.29
## Time         1.6622726      202467.06
```

Inference : Random Forest gives us the better model with the lowest RSquared error than the Random Tree. There is a problem of overfitting the data. Looks like Po2, Po1, NW, are very important predictors which contribute to the Random Trees by the measure of %IncMSE. Increasing the number of variables at each split decreases the accuracy of the model.

Question 10.2

Describe a situation or problem from your job, everyday life, current events, etc., for which a logistic regression model would be appropriate. List some (up to 5) predictors that you might use.

Answer:

Logistic Regression method can be used to predict the probability of success(pass/fail) in the exam based on the number of hours of study, class attendance of the student, homework submission, homework grade and student's health.

#####Q10.3#####

Question 10.3

1. Using the GermanCredit data set `germancredit.txt` from <http://archive.ics.uci.edu/ml/machine-learning-databases/statlog/german/> (description at <http://archive.ics.uci.edu/ml/datasets/Statlog+%28German+Credit+Data%29>), use logistic regression to find a good predictive model for whether credit applicants are good credit risks or not. Show your model (factors used and their coefficients), the software output, and the quality of fit. You can use the `glm` function in R. To get a logistic regression (logit) model on data where the response is either zero or one, use `family=binomial(link="logit")` in your `glm` function call.
2. Because the model gives a result between 0 and 1, it requires setting a threshold probability to separate between "good" and "bad" answers. In this data set, they estimate that incorrectly identifying a bad customer as good, is 5 times worse than incorrectly classifying a good customer as bad. Determine a good threshold probability based on your model.

```
set.seed(10)

rm(list = ls())
germancredit <- read.table("germancredit.txt",header = FALSE)
str(germancredit)

## 'data.frame':    1000 obs. of  21 variables:
## $ V1 : chr  "A11" "A12" "A14" "A11" ...
```

```
## $ V2 : int 6 48 12 42 24 36 24 36 12 30 ...
## $ V3 : chr "A34" "A32" "A34" "A32" ...
## $ V4 : chr "A43" "A43" "A46" "A42" ...
## $ V5 : int 1169 5951 2096 7882 4870 9055 2835 6948 3059 5234 ...
## $ V6 : chr "A65" "A61" "A61" "A61" ...
## $ V7 : chr "A75" "A73" "A74" "A74" ...
## $ V8 : int 4 2 2 2 3 2 3 2 2 4 ...
## $ V9 : chr "A93" "A92" "A93" "A93" ...
## $ V10: chr "A101" "A101" "A101" "A103" ...
## $ V11: int 4 2 3 4 4 4 4 2 4 2 ...
## $ V12: chr "A121" "A121" "A121" "A122" ...
## $ V13: int 67 22 49 45 53 35 53 35 61 28 ...
## $ V14: chr "A143" "A143" "A143" "A143" ...
## $ V15: chr "A152" "A152" "A152" "A153" ...
## $ V16: int 2 1 1 1 2 1 1 1 1 2 ...
## $ V17: chr "A173" "A173" "A172" "A173" ...
## $ V18: int 1 1 2 2 2 2 1 1 1 1 ...
## $ V19: chr "A192" "A191" "A191" "A191" ...
## $ V20: chr "A201" "A201" "A201" "A201" ...
## $ V21: int 1 2 1 1 2 1 1 1 1 2 ...
```

Notice that we have many categorical predictors.

#Make the response variable binary in terms of 0 and 1.

```
germancredit$V21[germancredit$V21==1] <- 0
germancredit$V21[germancredit$V21==2] <- 1
```

```
head(germancredit)
```

```
##      V1 V2  V3  V4   V5  V6  V7 V8  V9  V10 V11  V12 V13  V14  V15 V16  V17
V18
## 1 A11  6 A34 A43 1169 A65 A75  4 A93 A101  4 A121  67 A143 A152  2 A173
1
## 2 A12 48 A32 A43 5951 A61 A73  2 A92 A101  2 A121  22 A143 A152  1 A173
1
## 3 A14 12 A34 A46 2096 A61 A74  2 A93 A101  3 A121  49 A143 A152  1 A172
2
## 4 A11 42 A32 A42 7882 A61 A74  2 A93 A103  4 A122  45 A143 A153  1 A173
2
## 5 A11 24 A33 A40 4870 A61 A73  3 A93 A101  4 A124  53 A143 A153  2 A173
2
## 6 A14 36 A32 A46 9055 A65 A73  2 A93 A101  4 A124  35 A143 A153  1 A172
2
##      V19  V20 V21
## 1 A192 A201  0
## 2 A191 A201  1
## 3 A191 A201  0
## 4 A191 A201  0
## 5 A191 A201  1
## 6 A192 A201  0
```

```
#split the data into training and testing sets.
```

```
germancredit_train <- germancredit[1:800,]
```

```
germancredit_test <- germancredit[801:1000,]
```

```
#create a logistic regression model
```

```
germancredit_model = glm(V21~., family=binomial(link = "logit"),  
                          data=germancredit_train)
```

```
summary(germancredit_model)
```

```
##
```

```
## Call:
```

```
## glm(formula = V21 ~ ., family = binomial(link = "logit"), data =  
germancredit_train)
```

```
##
```

```
## Deviance Residuals:
```

```
##      Min       1Q   Median       3Q      Max  
## -2.7373  -0.6979  -0.3604   0.6663   2.5591
```

```
##
```

```
## Coefficients:
```

```
##              Estimate Std. Error z value Pr(>|z|)  
## (Intercept)  3.978e-01  1.249e+00   0.318  0.750139  
## V1A12        -2.701e-01  2.437e-01  -1.108  0.267784  
## V1A13        -9.306e-01  4.018e-01  -2.316  0.020567 *  
## V1A14        -1.737e+00  2.686e-01  -6.465  1.02e-10 ***  
## V2           2.893e-02  1.042e-02   2.777  0.005479 **  
## V3A31         2.255e-01  6.289e-01   0.359  0.719936  
## V3A32        -7.639e-01  4.753e-01  -1.607  0.108022  
## V3A33        -9.172e-01  5.233e-01  -1.753  0.079627 .  
## V3A34        -1.487e+00  4.907e-01  -3.031  0.002440 **  
## V4A41        -1.832e+00  4.425e-01  -4.141  3.46e-05 ***  
## V4A410       -1.413e+00  8.263e-01  -1.710  0.087326 .  
## V4A42        -9.368e-01  2.990e-01  -3.134  0.001727 **  
## V4A43        -9.044e-01  2.799e-01  -3.230  0.001236 **  
## V4A44        -8.312e-01  8.946e-01  -0.929  0.352807  
## V4A45        -3.222e-01  6.092e-01  -0.529  0.596843  
## V4A46         1.688e-02  4.255e-01   0.040  0.968354  
## V4A48        -2.213e+00  1.219e+00  -1.816  0.069365 .  
## V4A49        -8.368e-01  3.850e-01  -2.173  0.029760 *  
## V5           1.138e-04  5.166e-05   2.202  0.027682 *  
## V6A62        -3.991e-01  3.182e-01  -1.254  0.209771  
## V6A63        -4.615e-01  4.762e-01  -0.969  0.332404  
## V6A64        -1.222e+00  5.473e-01  -2.232  0.025592 *  
## V6A65        -7.093e-01  2.929e-01  -2.421  0.015462 *  
## V7A72        -2.017e-01  4.948e-01  -0.408  0.683485  
## V7A73        -3.028e-01  4.706e-01  -0.643  0.519975  
## V7A74        -1.105e+00  5.113e-01  -2.162  0.030623 *  
## V7A75        -4.092e-01  4.712e-01  -0.869  0.385102  
## V8           3.602e-01  9.933e-02   3.626  0.000287 ***  
## V9A92        -4.434e-01  4.300e-01  -1.031  0.302374
```

```
## V9A93      -1.230e+00  4.245e-01  -2.897 0.003769 **
## V9A94      -4.630e-01  5.119e-01  -0.905 0.365705
## V10A102     7.521e-01  4.771e-01   1.576 0.114917
## V10A103    -9.329e-01  4.830e-01  -1.931 0.053423 .
## V11         3.282e-03  9.850e-02   0.033 0.973420
## V12A122     4.101e-01  2.897e-01   1.415 0.156969
## V12A123     1.536e-01  2.649e-01   0.580 0.562115
## V12A124     7.122e-01  4.714e-01   1.511 0.130827
## V13        -1.868e-02  1.055e-02  -1.770 0.076682 .
## V14A142    -1.442e-02  4.733e-01  -0.030 0.975695
## V14A143    -4.354e-01  2.724e-01  -1.599 0.109919
## V15A152    -3.967e-01  2.739e-01  -1.448 0.147576
## V15A153    -5.576e-01  5.303e-01  -1.051 0.293071
## V16         3.297e-01  2.124e-01   1.552 0.120602
## V17A172     5.151e-01  7.807e-01   0.660 0.509351
## V17A173     5.655e-01  7.507e-01   0.753 0.451267
## V17A174     8.202e-01  7.597e-01   1.080 0.280307
## V18         5.065e-01  2.854e-01   1.775 0.075972 .
## V19A192    -3.739e-01  2.323e-01  -1.610 0.107489
## V20A202    -1.498e+00  8.079e-01  -1.854 0.063779 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 975.68  on 799  degrees of freedom
## Residual deviance: 705.07  on 751  degrees of freedom
## AIC: 803.07
##
## Number of Fisher Scoring iterations: 5

#consider doing some type of variable selection even though this has not
#been covered in the lectures yet. Also, notice how glm() implicitly
#creates dummy binary variables for each of the categorical variables.
#This is the correct way to do regression with categorical variables.
#however, if you want to do variable selection with these many dummy
#variables, you must re-define your categorical variables either manually
# or with an R function.

yhat<-predict(germancredit_model,germancredit_test[, -21],type= "response")
table(germancredit_test$V21, round(yhat))

##
##      0      1
## 0 115    24
## 1   29    32

#Important to use type = "response" here because without this
# we are given predictions of log-odds in the default case.
```

*#"round" your the yhat to get binary predictions from which
#you can compute an accuracy (classification rate). You may want
#to try out differnt thresholds for rounding. You can also use AUC to
#estimate the quality of fit.*

```
library(pROC)
```

```
## Type 'citation("pROC")' for a citation.
```

```
##
```

```
## Attaching package: 'pROC'
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
##      cov, smooth, var
```

```
roc(germancredit_test$V21, round(yhat))
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

```
##
```

```
## Call:
```

```
## roc.default(response = germancredit_test$V21, predictor = round(yhat))
```

```
##
```

```
## Data: round(yhat) in 139 controls (germancredit_test$V21 0) < 61 cases  
(germancredit_test$V21 1).
```

```
## Area under the curve: 0.676
```

*#Look at different threshold probability values and then compute the
#cost that corresponds to each threshold.*

```
thresh <- 0.8
```

```
yhat_thresh <- as.integer(yhat > thresh)
```

```
conf_matrix <- as.matrix(table(yhat_thresh, germancredit_test$V21))
```

```
conf_matrix
```

```
##
```

```
## yhat_thresh    0    1
```

```
##           0 134  53
```

```
##           1   5   8
```

```
accuracy <-
```

```
(conf_matrix[1,1]+conf_matrix[2,2])/(conf_matrix[1,1]+conf_matrix[1,2]+conf_m  
atrix[2,1]+conf_matrix[2,2])
```

```
accuracy
```

```
## [1] 0.71
```

```
specificity <- (conf_matrix[1,1])/(conf_matrix[1,1]+conf_matrix[2,1])
```

```
specificity
```

```
## [1] 0.9640288
```

```
thresh <- 0.7
```

```
yhat_thresh <- as.integer(yhat > thresh)
```

```
conf_matrix <- as.matrix(table(yhat_thresh, germancredit_test$V21))
```

```
conf_matrix
```

```
##
```

```
## yhat_thresh    0    1
```

```
##              0 132  43
```

```
##              1   7  18
```

```
accuracy <-
```

```
(conf_matrix[1,1]+conf_matrix[2,2])/(conf_matrix[1,1]+conf_matrix[1,2]+conf_m  
atrix[2,1]+conf_matrix[2,2])
```

```
accuracy
```

```
## [1] 0.75
```

```
specificity <- (conf_matrix[1,1])/(conf_matrix[1,1]+conf_matrix[2,1])
```

```
specificity
```

```
## [1] 0.9496403
```

Conclusion:

Based on the confusion Matrix for threshold 0.8 has accuracy 0.71 and specificity = 0.9640288 and for threshold 0.7 accuracy is 0.75 and specificity = 0.9496403. With threshold 0.7 we are getting better accuracy and specificity.